


Getting started with Caffe

Jon Barker, Solutions Architect



Agenda

- Caffe tour
 - Overview
 - Example applications
 - Setup
 - Performance
- Hands-on lab preview

A tour of Caffe

What is Caffe?

An open framework for deep learning developed by the Berkeley Vision and Learning Center (BVLC)

- Pure C++/CUDA architecture
- Command line, Python, MATLAB interfaces
- Fast, well-tested code
- Pre-processing and deployment tools, reference models and examples
- Image data management
- Seamless GPU acceleration
- Large community of contributors to the open-source project



caffe.berkeleyvision.org

<http://github.com/BVLC/caffe>

What is Caffe?

End-to-end Deep Learning for the practitioner and developer



Prototype



Train



Deploy

Caffe features

Data pre-processing and management

Data ingest formats

LevelDB or LMDB database

In-memory (C++ and Python only)

HDF5

Image files

Pre-processing tools

LevelDB/LMDB creation from raw images

Training and validation set creation with shuffling

Mean-image generation

Data transformations

Image cropping, resizing, scaling and mirroring

Mean subtraction

`$CAFFE_ROOT/build/tools`

Caffe features

Deep Learning model definition

- ▶ Protobuf model format
 - ▶ Strongly typed format
 - ▶ Human readable
 - ▶ Auto-generates and checks Caffe code
 - ▶ Developed by Google
 - ▶ Used to define network architecture and training parameters
 - ▶ No coding required!

```
name: "conv1"
type: "Convolution"
bottom: "data"
top: "conv1"
convolution_param {
    num_output: 20
    kernel_size: 5
    stride: 1
    weight_filler {
        type: "xavier"
    }
}
```

Caffe features

Deep Learning model definition

- ▶ Loss functions:
 - ▶ Classification
 - ▶ Softmax
 - ▶ Hinge loss
 - ▶ Linear regression
 - ▶ Euclidean loss
 - ▶ Attributes/multiclassification
 - ▶ Sigmoid cross entropy loss
 - ▶ and more...
- ▶ Available layer types:
 - ▶ Convolution
 - ▶ Pooling
 - ▶ Normalization
- ▶ Activation functions:
 - ▶ ReLU
 - ▶ Sigmoid
 - ▶ Tanh
 - ▶ and more...

Caffe features

Deep Neural Network training

Network training also requires no coding - just define a “solver” file

```
net: "lenet_train.prototxt"  
base_lr: 0.01  
momentum: 0.9  
max_iter: 10000  
snapshot_prefix: "lenet_snapshot"  
solver_mode: GPU
```

```
> caffe train --solver lenet_solver.prototxt --gpu 0
```

All you need to run
things on the GPU

Multiple optimization algorithms available: SGD (+momentum), ADAGRAD, NAG

Caffe features

Monitoring the training process

Output to stdout:

```
I0814 14:44:33.410693 2026435328 solver.cpp:294] Iteration 0, Testing net (#0)
I0814 14:44:35.697690 2026435328 solver.cpp:343]      Test net output #0: accuracy = 0.0931
I0814 14:44:35.697720 2026435328 solver.cpp:343]      Test net output #1: loss = 2.30247 (* 1 = 2.30247 loss)
I0814 14:44:35.718361 2026435328 solver.cpp:214] Iteration 0, loss = 2.30184
I0814 14:44:35.718392 2026435328 solver.cpp:229]      Train net output #0: loss = 2.30184 (* 1 = 2.30184 loss)
I0814 14:44:35.718400 2026435328 solver.cpp:486] Iteration 0, lr = 0.001
I0814 14:44:41.550972 2026435328 solver.cpp:214] Iteration 100, loss = 1.72121
I0814 14:44:41.550999 2026435328 solver.cpp:229]      Train net output #0: loss = 1.72121 (* 1 = 1.72121 loss)
I0814 14:44:41.551007 2026435328 solver.cpp:486] Iteration 100, lr = 0.001
I0814 14:44:47.383386 2026435328 solver.cpp:214] Iteration 200, loss = 1.73216
I0814 14:44:47.383415 2026435328 solver.cpp:229]      Train net output #0: loss = 1.73216 (* 1 = 1.73216 loss)
I0814 14:44:47.383424 2026435328 solver.cpp:486] Iteration 200, lr = 0.001
I0814 14:44:53.220012 2026435328 solver.cpp:214] Iteration 300, loss = 1.30751
I0814 14:44:53.220772 2026435328 solver.cpp:229]      Train net output #0: loss = 1.30751 (* 1 = 1.30751 loss)
I0814 14:44:53.220782 2026435328 solver.cpp:486] Iteration 300, lr = 0.001
I0814 14:44:59.053917 2026435328 solver.cpp:214] Iteration 400, loss = 1.16627
I0814 14:44:59.053948 2026435328 solver.cpp:229]      Train net output #0: loss = 1.16627 (* 1 = 1.16627 loss)
I0814 14:44:59.053956 2026435328 solver.cpp:486] Iteration 400, lr = 0.001
I0814 14:45:04.833677 2026435328 solver.cpp:294] Iteration 500, Testing net (#0)
I0814 14:45:06.778378 2026435328 solver.cpp:343]      Test net output #0: accuracy = 0.5589
I0814 14:45:06.778411 2026435328 solver.cpp:343]      Test net output #1: loss = 1.2699 (* 1 = 1.2699 loss)
```

To visualize - pipe, parse and plot or use DIGITS

Caffe features

Deep Neural Network deployment

Standard, compact model format

`caffe train` produces a binary `.caffemodel` file

Easily integrate trained models into data pipelines

Deploy against new data using command line, Python or MATLAB interfaces

Deploy models across HW and OS environments

`.caffemodel` files transfer to any other Caffe installation (including DIGITS)

Caffe features

Deep Neural Network sharing

Caffe Model Zoo hosts community shared models

Benefit from networks that you could not practically train yourself


<https://github.com/BVLC/caffe/wiki/Model-Zoo>

Caffe comes with unrestricted use of BVLC models:

AlexNet

R-CNN

GoogLeNet



Caffe features

Extensible code

Layer Protocol == Class Interface

```
import caffe
import numpy as np

class EuclideanLoss(caffe.Layer):

    def setup(self, bottom, top):
        # check input pair
        if len(bottom) != 2:
            raise Exception("Need two inputs to compute distance.")

    def reshape(self, bottom, top):
        # check input dimensions match
        if bottom[0].count != bottom[1].count:
            raise Exception("Inputs must have the same dimension.")
        # difference is shape of inputs
        self.diff = np.zeros_like(bottom[0].data, dtype=np.float32)
        # loss output is scalar
        top[0].reshape(1)

    def forward(self, bottom, top):
        self.diff[...] = bottom[0].data - bottom[1].data
        top[0].data[...] = np.sum(self.diff**2) / bottom[0].num / 2.

    def backward(self, top, propagate_down, bottom):
        for i in range(2):
            if not propagate_down[i]:
                continue
            if i == 0:
                sign = 1
            else:
                sign = -1
            bottom[i].diff[...] = sign * self.diff / bottom[i].num
```

Define a class in C++ or Python to extend Layer
Include your new layer in a network prototxt

```
layer {
    type: "Python"
    python_param {
        module: "layers"
        layer: "EuclideanLoss"
    }
}
```

Caffe example applications

Example applications

Use case 1: classification of images

Object

<http://demo.caffe.berkeleyvision.org/>

Open source demo code:

`$CAFFE_ROOT/examples/web_demo`

Scene

<http://places.csail.mit.edu/>

B. Zhou et al. NIPS 14

Style

<http://demo.vislab.berkeleyvision.org/>

Karayev et al. *Recognizing Image Style*. BMVC14




Maximally accurate	Maximally specific
cat	1.80727
domestic cat	1.74727
feline	1.72787
tabby	0.99133
domestic animal	0.78542



Example applications

Use case 2: localization




(Fast) Region based Convolutional Networks (R-CNN)

Ross Girshick, Microsoft Research

<https://github.com/rbgirshick/fast-rcnn>

Example applications

Use case 3: pixel level classification and segmentation




<http://fcn.berkeleyvision.org>

Long, Shelhamer, Darrell, *Fully convolutional networks for semantic segmentation*, CVPR 2015

Example applications

Use case 4: Sequence learning

- ▶ Recurrent Neural Networks (RNNs) and Long Short Term Memory (LSTM)
 - ▶ Video
 - ▶ Language
 - ▶ Dynamic data
- ▶ Current Caffe pull request to add support
 - ▶ <https://github.com/BVLC/caffe/pull/1873>
 - ▶ <http://arxiv.org/abs/1411.4389>




A group of young men playing a game of soccer.

Jeff Donahue et al.

Example applications

Use case 5: Transfer learning

Lots of data



Object Classifier



Dog vs. Cat

Top 10 in 10 mins after finetuning

- ▶ Just change a few lines in the model prototxt file

```
layer {  
    name: "data"  
    type: "Data"  
    data_param {  
        source: "ilsvrc12 train"  
        ...  
    }  
    ...  
}  
...  
layer {  
    name: "fc8"  
    type: "InnerProduct"  
    inner_product_param {  
        num_output: 1000  
        ...  
    }  
    ...  
}  
...  
layer {  
    name: "data"  
    type: "Data"  
    data_param: {  
        source: "dogcat train"  
        ...  
    }  
    ...  
}  
...  
layer {  
    name: "fc8-dogcat"  
    type: "InnerProduct"  
    inner_product_param {  
        num_output: 2  
        ...  
    }  
    ...  
}
```

Caffe setup and performance

Caffe setup

NVIDIA fork enables multiGPU: <https://github.com/NVIDIA/caffe>

- Tried and tested by BVLC on Ubuntu 14.04/12.04 and OS X 10.8+
- Also demonstrated to compile on RHEL, Fedora and CentOS
- Download source from <https://github.com/BVLC/caffe>
- Unofficial 64-bit Windows port <https://github.com/niuzhiheng/caffe>
- Linux setup (see <http://caffe.berkeleyvision.org/installation.html>)
 - Download
 - Install pre-requisites
 - Install CUDA and cuDNN for GPU acceleration
 - Compile using make


GPU acceleration

`-gpu N` flag tells `caffe` which gpu to use

Alternatively, specify `solver_mode: GPU` in `solver.prototxt`



cuDNN 4 Speedup vs. Non-Tiled FFT
VGG + Caffe, Fwd Pass Convolutional Layers
GeForce TITAN X, Core i7-4930 on Ubuntu 14.04 LTS





cuDNN 4 Speedup vs. CPU-only for batch = 1
AlexNet + Caffe, Fwd Pass Convolutional Layers
GeForce TITAN X, Core i7-4930 on Ubuntu 14.04 LTS

cuDNN integration

<http://developer.nvidia.com/cudnn>


Drop-in support

Install cuDNN, uncomment USE_CUDNN :=1 in `Makefile.config` before build



cuDNN 3 vs cuDNN 2 on Caffe, Ubuntu 14.04 LTS, Intel(R) Core(TM) i7-4930K CPU @ 3.40GHz, 24GB RAM, GeForce Titan X

Caffe model mobile deployment




▶ Jetson TX1

- ▶ NVIDIA Maxwell™ GPU with 256 NVIDIA® CUDA® Cores
- ▶ 4 GB LPDDR4 Memory, 16 GB eMMC 5.1 Flash Storage
- ▶ Connects to 802.11ac Wi-Fi and Bluetooth enabled devices10/100/1000BASE-T
- ▶ No need to change code
- ▶ Simply compile Caffe and copy a trained .caffemodel to TK1

*Source: <http://petewarden.com/2014/10/25/how-to-run-the-caffe-deep-learning-vision-library-on-nvidias-jetson-mobile-gpu-board/>

Hands-on lab preview

bit.ly/dlnvlab3



- ▶ Use data pre-processing tools
- ▶ Edit a network definition
- ▶ Train a model
- ▶ Improve classification accuracy by modifying network parameters
- ▶ Visualize trained network weights
- ▶ Deploy a model using Python

Deep Learning Lab Series Schedule

developer.nvidia.com/deep-learning-courses

- Review the other seminars in series

Seminar #2 - Introduction to DIGITs

Seminar #4 - Getting Started with the Theano Framework

Seminar #5 - Getting Started with the Torch Framework

Hands-on Lab

1. Create an account at nvidia.qwiklab.com
2. Go to “Getting started with Caffe” lab at bit.ly/dlnvlab3
3. Start the lab and enjoy!

Only requires a supported browser, no NVIDIA GPU necessary!

Lab is free until end of Deep Learning Lab series

